

EE323 Project 4

STCP: Implementing a Reliable Transport Layer

June 1, 2020

TA: Soobin Lee, Mangi Cho

EE323TA@gmail.com

STCP: a Simple Reliable Transport Layer

- Design and implement your own socket layer, **MYSOCK**, which supports reliable transport layer
 - Socket is a set of layers
 - You should implement only Transport layer, others are given.
- **STCP (Simple TCP)**
 - Reliable, connection-oriented, in-order, full duplex end-to-end delivery mechanism
 - Compatible with TCP (but, it is **NOT** TCP)
 - No flow control, No retransmission
 - Please refer to the provided specification when in doubt.

Milestones

- Make the client and server program work in a **reliable** network
- Reliable mode: NO packet drop or out-of-order delivery in the network
- Should meet all remaining functionalities to transmit packets correctly

Getting Started

- Read the **KLMS PDF file** and **RFC 793** carefully
- Download the STCP tarball from the KLMS and extract it on one of the lab machines
 - `$ tar xzvf assignment4.tar.gz`
- Check any compile errors with current Makefile
 - It should compile and run on any lab machines
 - The server and client will compile, but NOT run: They are dummy

Code Structure

Application Layer

- Simple, dummy client / server
 1. Client sends a request for a file
 2. Server transmits a file to client
 3. Client saves the transmitted file locally with filename 'recv'
- Help with debugging of your transport layer

Transport Layer: your task!

- Currently, it is a just bogus minimal transport layer
- You should implement your own transport layer in "transport.c"

Network Layer

- Emulates an unreliable datagram communication service
- Reliable or ~~unreliable~~ modes with client / server option -U
 - Default : reliable mode
 - ~~-U : unreliable mode~~
- Interfaces for transport layer is defined in stcp_api.h

mysock.h

stcp_api.h

MYSOCK Layer Overview

**Warning: this ppt does not cover the complete specification.
Please read the assignment material carefully.**

- Important implementation items
 - 3-way handshaking for connection establishment
 - Sequence number semantics for packets
 - Sliding windows for receiver and sender windows
Please check the rules for managing the windows
 - Slow start
 - 4-way handshaking for connection teardown
 - NO TCP option handling
 - No optimization for small data:
Send data packets as soon as data is available from the application
 - No delayed ACK:
Send ACK packets as soon as data is received

Your Working Playground Is ...

- **transport.c**
 - You will implement your transport layer in transport.c
 - You are **NOT** allowed to modify any other .c or .h files
 - You are **NOT** allowed to modify Makefile
 - Read comments in the file carefully to understand what to do
 - Consider **error or corner cases** and make sure to **clean up dynamically-allocated memory**
- One thread manages only one connection.
- Mysock.c calls transport_init() to make thread

extern void transport_init(mysocket_t sd, bool_t is_active);

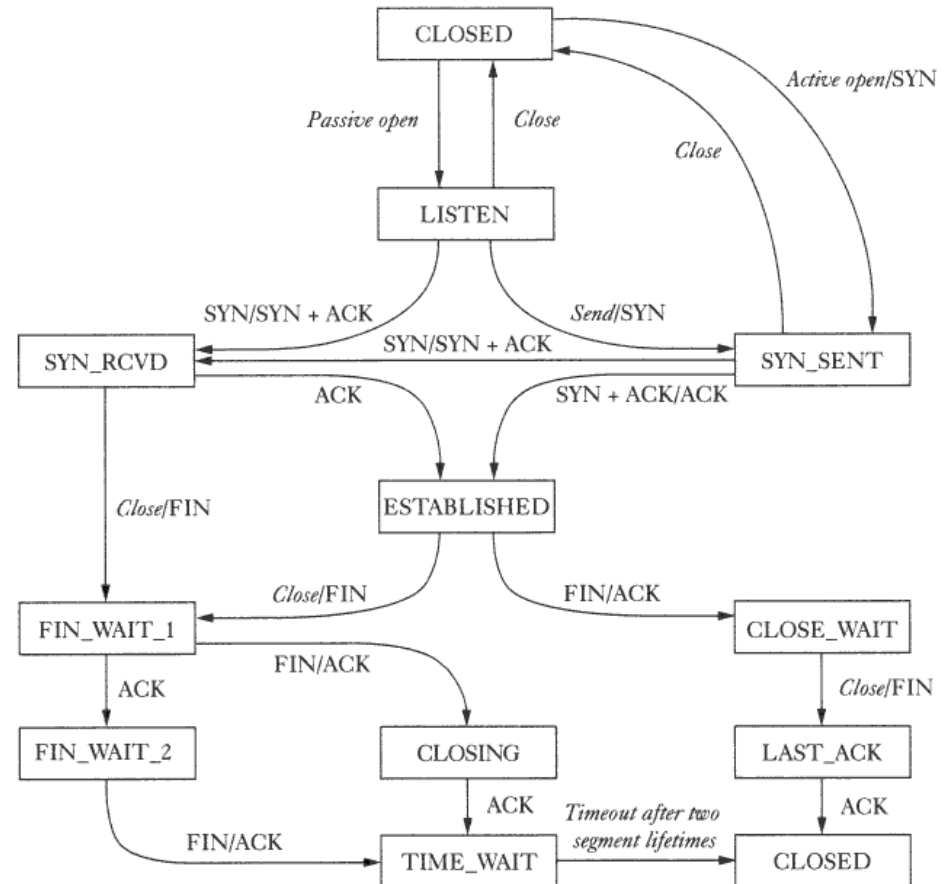
- A connection is initialized by calling transport_init();
- transport_init(); should not return until the connection ends

Implementation Hints

- Implement the TCP state diagram
- Need to have your **TCP context** (struct context_t) maintain a current state and react to the events from the application or the network layer possibly to transition from one state to other
- Include **header files** and **define** constants if you need
 - SEND_WINDOW_SIZE , RECV_WINDOW_SIZE, CWND_SIZE, PAYLOAD_SIZE, ...,
- Add **STCP states** at enum
 - CSTATE_CLOSED, CSTATE_LISTEN, CSTATE_SYN_SENT, CSTATE_SYN_RCVD, CSTATE_ESTABLISHED, CSTATE_FIN_WAIT_1, CSTATE_FIN_WAIT_2, CSTATE_CLOSE_WAIT, CSTATE_LAST_ACK, CSTATE_CLOSING
- Add **structures** if you need
 - STCPPacket, RxSegment, TxSegment, ...
- Create **functions** as freely as you want

TCP State Diagram

- Clearly understand the TCP state diagram by reading RFC 793
- What does each state mean?
- How transitions can be happened?
- What should be done in the current state?
- Check the differences between real TCP and our STCP



transport_init()

- All your STCP processing starts from here
- Make a **TCP context instance** and fill the **initial values**
- If **is_active == TRUE**, then your application wants to initiate a connection (e.g., called myconnect())
 - Create and send a **SYN segment** and mark your state to **SYN_SENT**
 - You may need to manage Tx packets to check ACK
- If **is_active == FALSE**, your application is listening on a port
 - Your TCP state should be **LISTEN**
- Call **control_loop()** for the main process

control_loop()

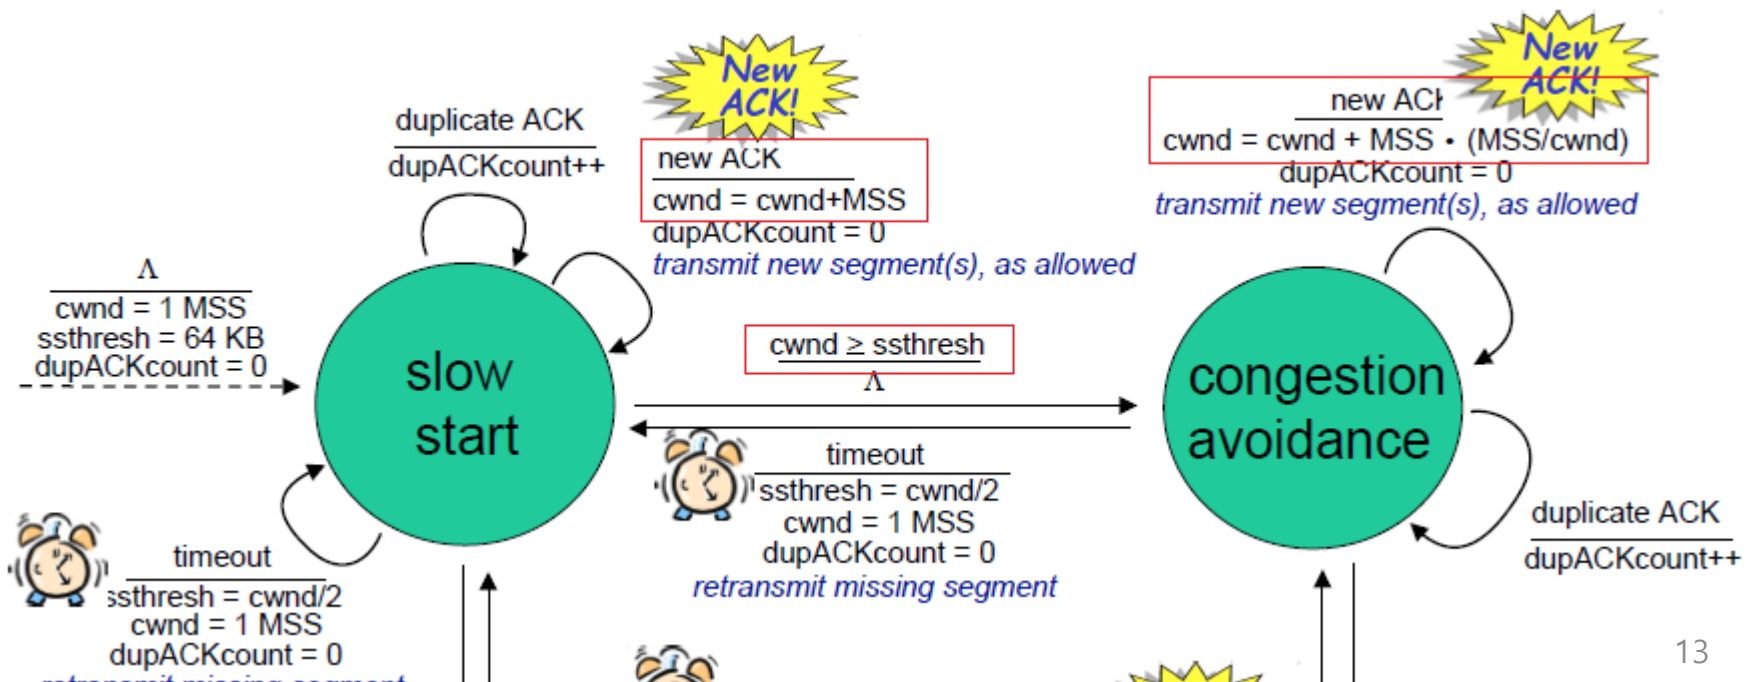
- **Main processes** are described in this function
- Get an **event** using **stcp_wait_for_event()**
 - Incoming data from the peer
 - New data from the application via mywrite()
 - The socket to be closed via myclose()
- Do appropriate jobs considering the event
 - Check the state, change the state, and send a packet, etc.
 - Close the connection
- Use **TCP context** to store the state and other necessary information
- Exit the control-loop when the connection is finished
- Use functions in **stcp_api.c** and **stcp_api.h**

stcp_api.c & stcp_api.h

- Network layer interface for transport layer is defined in stcp_api.h
- Important functions
 - stcp_unblock_application()
 - stcp_wait_for_event()
 - stcp_network_recv()
 - stcp_network_send()
 - stcp_app_recv()
 - stcp_app_send()
 - stcp_fin_received()
- See the descriptions on stcp_api.h for more details
- Highly recommend to study the implementation of these functions in network layer


Sliding window & Slow start

- You have to manage two window, cwnd, rwnd.
 - $swnd = \min(cwnd, rwnd)$
- cwnd increases with rule of slow start
 - Begin with 1MSS and threshold is 4 MSS



Print logs

- Whenever you send data or whenever you receive ACK, print log of (Use fopen, fprintf):
 - swnd : full sender window size
 - Rem : window size that usable but not yet sent in sender window
 - $Swnd - (NextSeqNum - SendBase)$
 - Byte : size of data which is sent or is ACKed.



- Format :
 - Make 2 log files "client_log.txt", "server_log.txt"
 - "Send:~~Wt~~swnd~~Wt~~Rem~~Wt~~Byte~~Wn~~"
 - Whenever you send data (not only ACK) -> stcp_network_send()
 - "Recv:~~Wt~~swnd~~Wt~~Rem~~Wt~~Byte~~Wn~~"
 - Whenever you receive ACK -> stcp_network_rcv(), th_flags & TH_ACK

Print logs (Hint)

- It can be confusing when you should print log exactly.
- Part of code structure might be :

```
If ( event & APP_DATA) {                /* there is something to send*/  
    if meet conditions for sending      /* window condition */  
    stcp_network_send();                /* send packet */  
    **print sending log here**  
    post-processing for sliding window ...  
}
```

```
If ( event & NETWORK_DATA) {            /* there is something to receive*/  
    stcp_network_rcv();                  /* receive packet */  
    if the packet is ACK  
    **print receiving log here**  
    post-processing for sliding window ...  
}
```


Testing Your STCP

- Run given **client** and **server** on different shells (2 PuTTY)
 - \$./server
 - Check the **port number** that is bound by server randomly
 - \$./client [SERVER_IP_ADDR]:[SERVER_PORT_NUM]
- Give **the name of the file** to client
 - You can use **-f [FILENAME]** option to give the name of the file to be transferred
 - \$./client [SERVER_IP_ADDR]:[SERVER_PORT_NUM] -f [filename]
 - You can give the name of the file at run-time
- Client may generate a request to server and server will transmit the file as a response
- The received file at client will be saved as **“rcvd”**
- Use **diff** command to compare the original file and **“rcvd”**

Testing Your STCP

- When you run client run-time (not use `-f` option), client and server don't call `myclose()` until forced quit.
- So, we will check log files only with `-f` option.
- Solution file also print logs properly only when `-f` option.
- We will provide additional code for check logics of log files.

Demo (server)



swnd Rem Byte

```
Send: 536 536 19
Send: 536 517 517
Recv: 536 0 19
Send: 1072 555 536
Recv: 1072 19 517
Send: 1608 1072 536
Send: 1608 536 536
Recv: 1608 0 536
Send: 2144 1072 536
Recv: 2144 536 536
Send: 2278 1206 536
Recv: 2278 670 536
Send: 2404 1332 536
Send: 2404 796 536
Send: 2404 260 260
Recv: 2404 0 536
Send: 2523 655 471
Recv: 2523 184 536
Send: 2636 833 536
Recv: 2636 297 536
Send: 2744 941 536
Recv: 2744 405 536
Send: 2848 1045 536
Recv: 2848 509 260
Send: 2948 869 536
Send: 2948 333 333
Recv: 2948 0 471
Send: 3045 568 536
Recv: 3045 32 536
Send: 3072 595 536
Recv: 3072 59 536
Send: 3072 595 536
```

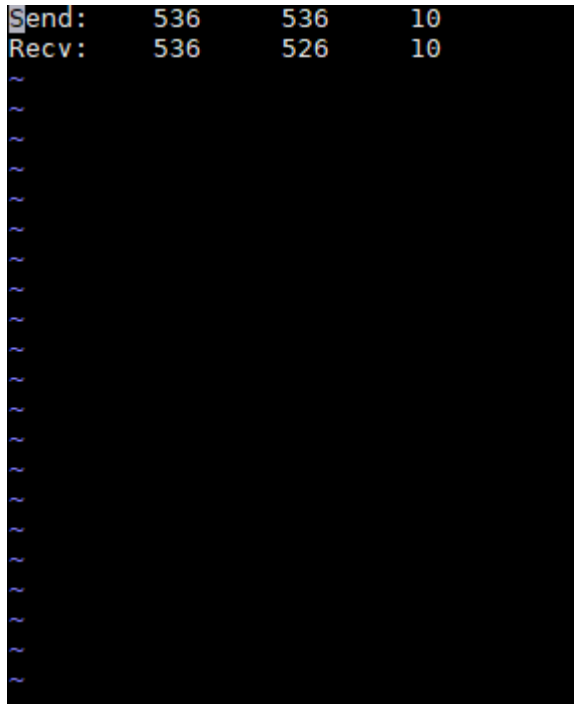
- Initially, swnd = 536, Rem = 536.
- Client request file to server, then Server send ACK to client
- (1 line) As response to request, server send data to client. At that time, swnd=536, Rem = 536, and sending data size = 19
- (2 line) There are more data to send, but the Rem is 517 (=536-19) so only send 517 byte.
- (3 line) ACK for first packet received (ACKed size 19), at that time Rem is 0. Then swnd increase by MSS (slow start)
- (4 line) Therefore, Rem is 555 (=536+19) so we can send data with maximum size 536.

The order can be different!
We will only check logic

Demo (client)

swnd Rem Byte

```
Send: 536 536 10
Recv: 536 526 10
```



- Simple!!
- Send request data (10 byte)
- Then received ACK. At that time Rem is 526
- If you have question about these demos, post piazza.

Tips

- **Print** every information that you want to check the correctness
 - Most straightforward and powerful way
- Do **NOT** try to implement everything at once
 - **Top-down implementation** is important
 - Implement **big branches** first
 - Just describe what should be done at each block briefly
 - Use **dummy function** that will be implemented later
 - Implement details **step-by-step** (test before implementing next block)
- Do **NOT** use global variable
 - Use **context instance** for each connection

Cautions

- You are NOT allowed to modify or submit any other .c , .h, or Makefile in stub codes rather than 'transport.c'.
 - You will submit only 'transport.c', but not other files.
 - You can modify code for your debugging, but remember that you code should work with original Makefile and supporting code.
- Your code will be graded on one of the lab machines.
 - Make sure that your code compiles and runs properly on the machines.
- We will test correct endianness.
 - Don't forgot to include your ntohs(), htons().

Submission

- Submission
 - transport.c: works in reliable mode
 - README: one page
 - Describing the design of your transport layer
 - Any design decisions/tradeoffs that you had to consider
 - Zipped like “YourStudentID_assign4.tar.gz”
 - `tar cvzf 20120000_assign4.tar.gz [source_files]`
 - Due: 11:59 pm on 6/26 , at KLMS

Others

- Do NOT copy and paste someone else's code including publicly available source code
- Start assignment *as quickly as possible.*
- *Design first, before you start it*
- This assignment is newly designed, so there might be some confusing points. If you have questions, feel free to post piazza.